

CHAPTER – 1

INTRODUCTION

1.1 Background

Artificial Intelligence (AI) has significantly transformed the way humans interact with computers. One of the most important applications of AI is the development of voice-based intelligent assistants, which enable users to communicate with computing systems using natural language. These assistants combine technologies such as speech recognition, Natural Language Processing (NLP), Machine Learning (ML), automation, and conversational AI to perform tasks efficiently and interactively.

Modern voice assistants such as Siri, Google Assistant, and Alexa demonstrate the potential of intelligent systems to automate daily activities and provide information through voice interaction. However, most existing assistants are primarily designed for mobile devices or cloud-based ecosystems, and they provide limited support for desktop-level automation and local processing.

The Jarvis AI — Desktop Based Voice Assistant project is developed to demonstrate how multiple AI technologies can be integrated into a single intelligent system capable of understanding voice commands, processing natural language, executing system-level tasks, and interacting conversationally with the user.

The project integrates several AI and software components, including:

- Speech-to-text processing (offline and online)
- Neural network-based intent classification
- Named entity recognition
- Sentiment and topic analysis
- Command parsing and automation
- Large Language Model (LLM) conversation handling
- Text-to-speech response generation
- Workflow automation and scheduling

By combining these technologies, Jarvis AI provides a natural, intelligent, and interactive desktop assistant experience.

1.2 Problem Statement

Despite the rapid advancement of voice assistants, many existing systems face several limitations.

Most assistants:

- are optimized for mobile platforms rather than desktop environments,
- rely heavily on cloud-based services,
- use rule-based command execution,
- lack context awareness and conversational intelligence,
- raise privacy concerns due to continuous internet dependency.

Desktop automation tools, on the other hand, often require manual interaction through keyboard and mouse, which can be inefficient for repetitive tasks.

Therefore, there is a need for a system that:

- understands natural voice commands,
- performs desktop automation tasks,
- works in both offline and online modes,
- integrates machine learning and NLP techniques,
- supports conversational interaction,
- maintains privacy through local processing.

The Jarvis AI project addresses these challenges by developing a hybrid AI assistant system for desktop environments.

1.3 Objectives

The primary objective of the Jarvis AI project is to design and implement a desktop-based intelligent voice assistant capable of understanding user commands and performing automation tasks.

Specific objectives include:

- **To develop a voice-controlled desktop assistant system** — Enables users to interact with the computer using natural voice commands.
- **To implement speech recognition for voice input processing** — Converts spoken audio into text for further processing.

- **To design a machine learning–based intent classification system** — Identifies the user’s intention from the recognized text.
- **To integrate Natural Language Processing components** — Improves understanding of language through entity recognition, sentiment detection, and topic analysis.
- **To implement command parsing and automation modules** — Breaks commands into executable tasks and performs system operations.
- **To integrate LLM-based conversational AI capabilities** — Allows the assistant to handle open-ended conversations intelligently.
- **To support task scheduling and workflow automation** — Enables execution of commands automatically at specified times or sequences.
- **To create a modular and extensible assistant architecture** — Ensures the system can be easily updated with new features.

1.4 Scope of the Project

The Jarvis AI system focuses on building a desktop-based intelligent assistant that improves user interaction with the computer through voice commands.

- **Voice command recognition** — Allows users to control the system using speech input.
- **Natural language understanding** — Enables the assistant to interpret user commands meaningfully.
- **Desktop application control** — Provides voice-based control for opening and managing applications.
- **Information retrieval** — Allows the assistant to fetch data such as weather, news, or search results.
- **Conversational interaction** — Supports natural communication between the user and the assistant.
- **Sentiment and topic analysis** — Helps analyse and classify user input text.
- **Translation support** — Enables multilingual communication through translation features.
- **Workflow automation** — Allows execution of multiple commands as a single routine.
- **Hybrid offline–online operation** — Ensures the assistant works with or without internet connectivity.

1.5 Organization of the Report

The report is organized into the following chapters:

- **Chapter 1 — Introduction:** Overview of the project, objectives, and scope.
- **Chapter 2 — Literature Review:** Study of existing voice assistant systems and NLP techniques.
- **Chapter 3 — System Analysis:** Existing system analysis, proposed system, feasibility study, and requirements.
- **Chapter 4 — System Design:** Architecture design, DFD, UML diagrams.
- **Chapter 5 — Implementation:** Tools, technologies, module implementation, and algorithms.
- **Chapter 6 — Testing:** Testing methods, test cases, and validation.
- **Chapter 7 — Results and Discussion:** Output analysis and system performance.
- **Chapter 8 — Conclusion and Future Scope:** Summary and possible improvements.
- **Chapter 9—References:** Websites and research papers referred.
- **Chapter 10—Appendices:** Sample outputs.

CHAPTER – 2

LITERATURE REVIEW

2.1 Introduction to Voice Assistant Systems

Voice assistants represent one of the most practical applications of Artificial Intelligence in human–computer interaction. These systems integrate speech recognition, natural language processing, machine learning, and automation technologies to allow users to interact with computing systems through voice commands.

Recent research in AI-based assistants has focused on improving:

- intent recognition accuracy,
- automation capabilities,
- and privacy-aware processing.

Most early assistant systems were rule-based, while modern systems increasingly rely on machine learning and deep learning models for language understanding.

2.2 Study of Existing Voice Assistant Systems

Several research works and implementations have contributed to the development of voice-based intelligent systems.

Paper 1: AI Powered Voice Assistant for Enhanced User Interaction

This paper analyses AI voice assistants with personalization and task-management features, focusing on improving user interaction through voice customization, NLP-based task execution, and multi-platform accessibility. The system integrates speech recognition, scheduling, information retrieval, and personalization to create an intelligent virtual assistant experience [1].

Paper 2: Voice Assistant with AI Chat Integration using OpenAI

This research integrates a voice assistant with OpenAI conversational AI models to enable natural, context-aware conversations. The system combines speech recognition, NLP, and GPT-based chat capabilities using a web interface to provide intelligent responses, personalized assistance, and improved accessibility [2].

Paper 3: Embedded End-to-End Voice Assistant

This system explores offline automatic speech recognition (ASR), deep learning models, and embedded hardware integration to create a self-contained assistant environment. The design emphasizes local processing and reduced network dependency [3].

Paper 4: Automate Personal Voice Based Assistant Using Python

This paper develops a Python desktop voice assistant that performs tasks such as setting reminders, opening applications, playing music, and retrieving information using voice commands. The system demonstrates the integration of speech recognition, NLP processing, and automation features in a desktop environment [4].

Paper 5: AI Voice Assistant using NLP and Python Libraries (M.A.R.F)

This paper presents a Python-based voice assistant chatbot (M.A.R.F.) designed to automate daily tasks such as playing music, retrieving information, opening applications, and web searching. The system integrates speech recognition, automation features, multilingual support, and memory storage capabilities to improve productivity and interaction [5].

2.3 Natural Language Processing Techniques in Voice Assistants

Modern voice assistants rely on Natural Language Processing (NLP) techniques to understand and interpret user commands effectively.

Key NLP techniques include:

- **Tokenization** — Splitting user input into smaller words or tokens.
- **Lemmatization** — Converting words into their base or root form.
- **Intent Classification** — Identifying the purpose of a user command.
- **Named Entity Recognition (NER)** — Extracting important entities such as names or locations.
- **Sentiment Analysis** — Determining the emotional tone of text.
- **Topic Classification** — Categorizing text into domains like technology or finance.

These NLP components enable the assistant to understand natural language intelligently instead of relying only on predefined commands.

2.4 Speech Recognition Technologies

Speech recognition is a core component of voice assistants that converts spoken audio into text.

Two main approaches are used:

1) Offline Speech Recognition

- Processes audio locally on the system
- Improves privacy and reliability

2) Online Speech Recognition

- Uses cloud-based speech processing services
- Provides higher transcription accuracy

Combining both methods creates a hybrid speech recognition system.

2.5 Conversational AI and Large Language Models

Conversational AI systems use Large Language Models (LLMs) to generate intelligent responses for user queries.

This approach combines:

- Rule-based automation
- Machine learning intent classification
- Generative AI conversation

This enables more natural and intelligent interaction between the user and the assistant.

2.6 Comparative Analysis of Reviewed Papers

Paper Name	Technologies Used	Limitations
An AI Powered Voice Assistant for Enhanced User Interaction	Speech Recognition, NLP (spaCy / NLTK), Python libraries (pyttsx3), TTS	Primarily conceptual lacks deep contextual understanding
Voice Assistant with AI Chat Integration using OpenAI	OpenAI GPT models, Speech-to-Text, NLP, Cloud APIs	Requires internet connectivity for AI model access.
An Embedded End-to-End Voice Assistant	Deep Learning, Offline ASR, Transfer Learning, Embedded Systems, Edge computing, Mozilla DeepSpeech / NeMo toolkit, Rasa / spaCy (NLU)	Requires high-end embedded hardware focused mainly on automotive environments
Automate Personal Voice Based Assistant Using Python	Python, SpeechRecognition, pyttsx3, rule-based NLP, CustomTkinter, PyWhatKit	Mostly library-based implementation rather than ML-driven intelligence.
AI Voice Assistant using NLP and Python Libraries (M.A.R.F)	NLP, pyttsx3, web scraping (BeautifulSoup), pywhatkit, Google Translate API	Primarily rule-based task execution

Table 1.0 – Literature Review Analysis

2.7 Summary of Literature Review

From the study of existing systems, the following observations can be made:

- Many assistants rely on rule-based command execution.
- Cloud-dependent assistants raise privacy concerns.
- Offline assistants often require specialized hardware.
- Several systems lack context awareness and conversational intelligence.
- Few systems focus specifically on desktop automation.

These research gaps motivate the development of Jarvis AI — Desktop Based Voice Assistant, which integrates speech recognition, NLP, machine learning, and automation into a modular assistant system.

CHAPTER – 3

SYSTEM ANALYSIS

3.1 Existing System

Voice assistant systems have become an important application of Artificial Intelligence, enabling users to interact with computing devices using natural language. Existing voice assistants typically integrate speech recognition, natural language processing, and automation features to perform tasks such as searching the web, opening applications, and answering user queries.

Many assistant systems are implemented using speech recognition libraries, rule-based command matching, and text-to-speech systems. These assistants follow a sequential workflow in which speech input is converted to text, processed using predefined rules, and mapped to system actions.

While such systems improve usability and productivity, they often suffer from several limitations. Most assistants rely heavily on predefined command patterns, which restrict flexibility and reduce the ability to interpret natural language variations. Additionally, many modern assistants depend on cloud-based services for speech recognition and conversational intelligence. This continuous dependency on the internet raises privacy concerns, latency issues, and reliability problems.

Some research-based assistant systems support offline functionality, but they often require high-performance hardware or embedded environments, which reduces practicality for general desktop users.

Another limitation of existing systems is the lack of context awareness and conversational intelligence, making interaction less natural and adaptive. Many assistants cannot learn from user behaviour or handle compound commands effectively.

Because of these limitations, there is a need for a desktop-based intelligent assistant system that integrates machine learning, NLP, automation, and conversational AI in a modular and privacy-aware architecture.

3.2 Proposed System

The proposed system, Jarvis AI — Desktop Based Voice Assistant, is designed to provide an intelligent and flexible solution for desktop automation through natural voice interaction.

The system integrates multiple technologies, including:

- Speech recognition
- Natural Language Processing
- Machine Learning–based intent classification
- Automation modules
- Conversational AI
- Text-to-speech response generation

The proposed assistant follows a hybrid architecture, allowing it to operate in both offline and online modes. Offline processing ensures privacy and reliability, while online services provide advanced conversational capabilities.

The system workflow is as follows:

1. The user provides a voice command through a microphone.
2. The speech recognition module converts audio into text.
3. The NLP module processes the text and identifies the user’s intent.
4. The command parser extracts relevant information.
5. The automation module executes the requested task.
6. The assistant generates a response using text or speech output.

Unlike rule-based assistants, Jarvis AI uses intent-based command execution, which improves flexibility and natural interaction. The system is designed using a modular architecture, allowing new features and modules to be added easily.

The assistant can perform tasks such as:

- Opening applications
- Searching information online
- Managing workflows
- Translating text
- Conversational interaction
- Scheduling tasks
- Performing desktop automation

This makes the system suitable for personal productivity, accessibility support, and intelligent desktop interaction.

3.3 Feasibility Study

A feasibility study evaluates whether the proposed system can be successfully implemented from technical, economic, and operational perspectives.

Technical Feasibility

The Jarvis AI system is technically feasible because it is built using well-established programming tools and AI libraries.

The project uses:

- Python programming language
- TensorFlow/Keras
- NLTK
- TextBlob
- Vosk speech recognition
- Google Speech API
- Groq LLM API

These tools are widely used, well-documented, and suitable for building AI-based assistant systems. The system does not require specialized hardware and can run on a standard desktop or laptop with microphone support.

The modular architecture also ensures easy maintenance and scalability.

Economic Feasibility

The project is economically feasible because it uses open-source software tools and freely available libraries. Development can be completed using existing hardware resources without additional cost.

Optional cloud services such as LLM APIs can be used selectively, which helps minimize operational expenses.

Therefore, the overall cost of development and deployment is low.

Operational Feasibility

The system is designed to be user-friendly and practical for everyday use. Users interact with the assistant using simple voice commands, reducing the need for manual keyboard or mouse interaction.

The assistant automates repetitive tasks and improves productivity. The system can operate in both online and offline modes, ensuring reliability in different environments.

Because of its intuitive design and modular structure, the system is operationally feasible.

3.4 Requirements Analysis

Requirements analysis defines what the system must do and how it should perform.

1. Functional Requirements

The Jarvis AI system should be able to:

- Accept voice input from the user
- Convert speech to text
- Process natural language commands
- Classify user intent using machine learning
- Extract entities from user commands
- Execute desktop automation tasks
- Perform web searches
- Provide conversational responses
- Generate speech output
- Schedule tasks and workflows
- Store conversation history
- Generating images

2. Non-Functional Requirements

- **Performance:**

The system should respond to user commands with minimal delay.

- **Reliability:**

The assistant should function correctly in both offline and online modes.

- **Usability:**

The system should support simple and natural voice interaction.

- **Security and Privacy:**

Sensitive operations should be processed locally whenever possible.

- **Scalability:**

The modular design should allow easy addition of new features.

- **Maintainability:**

The system should be easy to update and debug.

CHAPTER – 4

SYSTEM DESIGN

4.1 Introduction

System design defines how the Jarvis AI assistant is structured and how different modules interact to process user commands. The design focuses on modularity, scalability, and efficient communication between components such as speech recognition, natural language processing, automation, and conversational AI.

The Jarvis AI system is designed using a layered architecture, where each layer performs a specific function in the voice-command processing pipeline. This approach simplifies development, testing, and maintenance.

4.2 System Architecture Design




Fig 1.1 (System Architecture Diagram)

The system architecture diagram illustrates the layered design of the Jarvis AI desktop voice assistant. The architecture begins with the Voice Interface, where the user interacts with the system using spoken commands. The speech input is processed by the Speech-to-Text (STT) module, which converts audio signals into text.

The text is then passed to the NLP Module, which performs language understanding tasks such as intent detection and context management. The Dialogue Manager and Knowledge Graph help maintain conversational context and provide intelligent responses.

The processed command is forwarded to the Task Execution Layer, which includes:

- System Control
- Web Services
- Automation Module

These components execute user commands such as opening applications, retrieving information, and performing automated workflows.

4.3 Data Flow Diagram (DFD)

Data Flow Diagrams illustrate how information moves through the Jarvis AI system.

4.3.1 Level-0 DFD




Fig 1.2 (Level 0 DFD)

The Level-0 DFD represents the system as a single process interacting with the user.

External Entity: User

Process: Jarvis AI System

Data Flows:

- User Commands
- Recognized Text
- Command Execution
- Processed Output

This diagram provides a high-level overview of the assistant system.

4.3.2 Level-1 DFD




Fig 1.3 (Level 1 DFD)

The Level-1 DFD provides a more detailed view of internal processes within the Jarvis AI system.

The system is divided into multiple functional modules:

- Voice Interface
- NLP Processing
- Task Execution
- Smart Home Control
- App Control

The user provides voice commands through the Voice Interface, which sends input to the NLP Processing module for intent recognition. The recognized intent is forwarded to the Task Execution module, which performs the requested action.

The system interacts with:

- External API services
- Application control modules

The database stores conversation data and system information, while external APIs provide additional services such as speech recognition and web search.

The system then generates responses and sends them back to the user.

4.4 UML Diagrams

UML diagrams are used to represent the structural and behavioural aspects of the Jarvis AI system.

4.4.1 Use Case Diagram




Fig 1.4 (Usecase Diagram)

The Use Case Diagram represents the interaction between the User and the Jarvis AI system. The primary actor is the user, who interacts with the assistant through voice commands. The diagram shows the major functionalities supported by the system, including:

- Giving voice commands
- Executing automation tasks
- Searching information
- Conversational interaction
- Translation
- Scheduling tasks
- Generating images

The diagram also introduces an Admin actor, who is responsible for managing and maintaining the assistant system. The Admin can perform system-level operations such as:

- Managing machine learning models
- Managing external APIs

While the User focuses on interaction and task execution, the Admin ensures that the system operates correctly and remains updated.

4.4.2 Sequence Diagram

The Sequence Diagram illustrates how modules interact during voice command execution.




Fig 1.5 (Sequence Diagram)

The sequence diagram illustrates the step-by-step interaction between different modules of the Jarvis AI system during authentication, hotword detection, and command execution.

The process begins when the user initiates authentication through the Face Authentication System. The system verifies the user's identity using the `verifyFace()` function. Once verification is successful, the system returns an authentication confirmation and allows further interaction with the assistant.

After authentication, the Hotword Service module starts listening for the activation phrase. When the hotword is detected, the system initializes the assistant and activates the speech processing modules. The Speech Module begins capturing audio input from the user.

The captured audio is passed to the Speech-to-Text (STT) Engine, which transcribes the voice input into text. The text is then forwarded to the NLP Engine, where it is processed to identify intent and extract relevant information from the command.

Next, the Task Router module determines which system component should handle the request. The request is then sent to the Automation Engine, which executes the required operation. If the command involves scheduling, the Scheduler module processes the request. For conversational queries, the system sends a request to the LLM API module, which generates an intelligent response.

Finally, the system returns the result to the user as output, completing the interaction cycle. This sequence diagram demonstrates how multiple modules collaborate in a structured workflow to process user commands from authentication to response generation.


Stepwise Working:

The sequence diagram represents the workflow of the Jarvis AI assistant as follows:

- The user initiates authentication through the Face Authentication System.
- The system verifies the user's identity using face recognition.
- After successful authentication, the hotword service starts listening.
- The user speaks the activation phrase (hotword).
- The system initializes assistant modules.
- The Speech Module captures voice input.
- The STT Engine converts speech into text.
- The NLP Engine processes the text and identifies intent.
- The Task Router determines the appropriate module.
- The Automation Engine executes the command.
- The Scheduler handles scheduled tasks when required.
- The LLM API generates responses for conversational queries.
- The system returns the response to the user.
- The assistant waits for the next command.

4.4.3 Class Diagram

The Class Diagram represents the internal structure of the system and relationships between modules.



The class diagram represents the structural design of the Jarvis AI assistant system. It shows the main classes in the system, their responsibilities, and the relationships between core system modules and service modules. The diagram is divided into two major sections: the Core System and Service Modules.

The Core System contains the primary classes responsible for voice processing, command understanding, authentication, and task execution. The central class is JarvisMain, which initializes the assistant and starts the listening process. The ListenModule captures audio input from the user, and the SpeechModule converts speech to text and generates spoken responses. The Brain class processes user commands and delegates tasks to appropriate modules.

The CommandParser class interprets processed commands and extracts relevant instructions, while the TaskExecutor executes those commands. The FaceAuthModule handles user authentication, ensuring secure access to the assistant. The AutomationEngine coordinates automation tasks and connects the core system to external service modules.

The Service Modules represent functional components that perform specific operations. These include modules for web search, weather information, email sending, WhatsApp messaging, file searching, translation, calculation, and news retrieval. Each service module contains a method responsible for performing its specific task, such as searchWeb(), getWeather(), or translateText().

The relationships shown in the diagram indicate how the core assistant system communicates with service modules through the automation engine. This modular class structure improves maintainability, scalability, and code organization by separating system logic from service-specific functionality.

Key Points

- The diagram shows the static structure of the Jarvis AI system.
- The system is divided into Core System classes and Service Modules.
- **JarvisMain** controls system initialization and listening.
- **ListenModule** captures user audio input.
- **SpeechModule** performs speech-to-text and text-to-speech operations.
- **Brain** processes commands and delegates tasks.
- **CommandParser** interprets commands.
- **TaskExecutor** executes actions.
- **FaceAuthModule** handles authentication.
- **AutomationEngine** connects the core system with service modules.
- Service modules implement features like search, translation, messaging, and news retrieval.
- The modular design improves scalability and maintainability.

4.5 Input Design and Output Design




Fig 1.7 (Input and Output Design Diagram)

4.5.1 Input Design

The system supports both voice input and text input.

Voice input:

Captured through a microphone and processed using speech recognition.

Text input:

Used in console mode or web interface.

Example inputs:

- Open Chrome
- Search Google for AI
- Translate hello to Hindi
- Set alarm for 7 AM

Input validation ensures reliable command processing.

4.5.2 Output Design

The assistant provides output in text and speech formats.

Text output:

Displayed in console or user interface.

Voice output:

Generated using text-to-speech modules.

Example outputs include:

- Spoken responses
- Automation confirmations
- Search results
- Conversational replies

The output design ensures smooth communication between the user and the assistant.

CHAPTER – 5

IMPLEMENTATION

5.1 Introduction

This chapter explains how the Jarvis AI — Desktop Based Voice Assistant system is implemented using Python and various AI technologies. The implementation integrates speech recognition, natural language processing, machine learning–based intent classification, automation, and conversational AI modules into a single assistant system.

The system follows a modular implementation approach where each component performs a specific function in the voice-command processing pipeline. This improves maintainability and scalability of the assistant.

The overall workflow of the implementation includes:

- Capturing voice input from the user
- Converting speech to text
- Processing text using NLP techniques
- Identifying user intent using a machine learning model
- Executing automation tasks
- Generating text and voice responses

Each of these stages is implemented as a separate module.

5.2 Tools and Technologies Used

The Jarvis AI system is implemented using several tools and libraries.

1) Programming

- Python

2) Machine Learning and NLP

- TensorFlow / Keras
- Hugging Face Transformers
- NLTK

- TextBlob

3) Speech Processing

- Vosk Speech Recognition (For offline speech processing)
- Google Speech Recognition API (For online speech processing)
- pyttsx3
- gTTS

4) Generative AI

- Groq LLM API (LLaMA models)

5) Computer Vision

- MediaPipe FaceMesh

6) Storage

- JSON files
- Pickle files
- Model files (.keras)

These technologies collectively enable intelligent assistant functionality.

5.3 Module Description

The Jarvis AI assistant is implemented using multiple functional modules. Each module performs a specific task in the assistant workflow.

1) Hotword Detection Module

The hotword detection module continuously listens for activation phrases such as “Hey Jarvis” or “Wake up.” This ensures that the assistant only activates when required, reducing unnecessary processing.

Key responsibilities include:

- Listening for activation keywords
- Activating the assistant system
- Switching to command-listening mode

This module improves system efficiency and usability.

2) Speech Recognition Module

The speech recognition module captures audio input from the user and converts it into text for further processing.

The system supports two recognition modes:

- **Offline Speech Recognition** — Implemented using Vosk models to process audio locally.
- **Online Speech Recognition** — Implemented using Google Speech Recognition API for higher transcription accuracy.

The recognized text is forwarded to the NLP module.

3) NLP Processing Module

The NLP module processes the recognized text to understand the user's command.

The module performs the following operations:

- **Tokenization** — Splits input text into individual words.
- **Lemmatization** — Converts words into their root forms.
- **Named Entity Recognition** — Extracts entities such as names or locations.
- **Sentiment Analysis** — Detects emotional tone.
- **Topic Classification** — Categorizes user input.

These steps transform raw text into structured information.

4) Intent Classification Module

Intent classification is implemented using a neural network trained on predefined command patterns.

The module performs the following steps:

- Convert text into Bag-of-Words representation
- Load trained model (jarvis_model.keras)
- Predict intent probabilities
- Select highest-probability intent

This approach improves command recognition accuracy.

5) Command Processing Module

The command processing module interprets classified intent and determines the corresponding action.

Key functions include:

- Extracting command parameters
- Handling compound commands
- Routing commands to automation module

This module connects NLP processing with system execution.

6) Automation Module

The automation module executes user commands on the system.

It supports operations such as:

- Opening applications
- Searching the web
- Scheduling tasks
- Running workflows
- Controlling system settings

This module acts as the execution engine of Jarvis AI.

7) Conversational AI Module

The conversational module handles queries that do not match predefined intents.

It performs the following tasks:

- Sending queries to Groq LLM API
- Generating intelligent responses

- Maintaining conversation flow

This allows the assistant to support natural conversation.

8) Text-to-Speech Module

The text-to-speech module converts system responses into spoken output.

Its functions include:

- Converting text responses into audio
- Playing audio through speakers
- Providing user feedback

This module completes the voice interaction cycle.

5.4 Algorithms / Logic

Voice Command Processing Algorithm

The assistant processes voice commands using the following steps:

- Face authentication
- Start hotword listener
- Detect activation command
- Capture voice input
- Convert speech to text
- Perform NLP preprocessing
- Classify user intent
- Execute task
- Generate response
- Convert response to speech
- Wait for next command

Intent Classification Logic

The intent classification model follows these steps:

- Load trained neural network
- Preprocess input text

- Convert text to Bag-of-Words vector
- Predict intent probabilities
- Return highest-probability intent

5.5 Code Snippets

This section presents selected code snippets from the project to highlight the core implementation logic of important modules. Instead of including the complete source code, only relevant portions are shown to demonstrate how key functionalities are designed and executed.

5.5.1 Intent Classification

File: [Brain.py]

```
# Brain.py - Neural intent predictor

import nltk

from nltk.stem import WordNetLemmatizer

from functools import lru_cache

import numpy as np

from tensorflow.keras.models import load_model

lemmatizer = WordNetLemmatizer()

_model = None

def _load_model():
    global _model
    if _model is None:
        import pickle, json
        _model = load_model('Models/jarvis_model.keras')
    return _model

def bag_of_words(sentence):
    """Convert to feature vector"""
    import pickle
    words = pickle.load(open('Models/words.pkl', 'rb'))
    tokens = nltk.word_tokenize(sentence.lower())
```



```
tokens = [lemmatizer.lemmatize(w) for w in tokens]
bag = [1 if w in tokens else 0 for w in words]
return np.array(bag)
```

5.5.2 Speech Recognition

File: [Listen.py]

```
# Listen.py - Async STT with noise filtering
import speech_recognition as sr
import noisereduce as nr
import webrtcvad
import numpy as np
from speech import speak
import webrtcvad
from deep_translator import GoogleTranslator

def apply_noise_cancellation(audio_data, sr_rate=16000):
    """VAD + Noise Reduction in 30ms frames"""
    vad = webrtcvad.Vad()
    vad.set_mode(1) # Balanced
    audio_int16 = (audio_data * 32768).astype(np.int16)
    frame_size = int(sr_rate * 0.03) # 30ms
    voiced = []
    for i in range(0, len(audio_int16), frame_size):
        frame = audio_int16[i:i+frame_size]
        if len(frame) < frame_size:
            frame = np.pad(frame, (0, frame_size - len(frame)))
            if vad.is_speech(frame.tobytes(), sr_rate):
                voiced.append(frame)
    if not voiced:
        return audio_data
    vad_audio = np.concatenate(voiced).astype(np.float32) / 32768.0
```

```
clean = nr.reduce_noise(y=vad_audio, sr=sr_rate, prop_decrease=0.3)
return clean

def listen():
    """Online STT with noise cancellation"""
    try:
        rec = sr.Recognizer()
        with sr.Microphone() as source:
            rec.adjust_for_ambient_noise(source, duration=1)
            audio = rec.listen(source, timeout=3, phrase_time_limit=8)
            query = rec.recognize_google(audio, language='en-in')
        return query.lower()
    except:
        return "None"
```

5.5.3 Image Generation

File: [ImageGeneration.py]

```
async def open_images(prompt):
    """Open generated images after creation"""
    folder_path = DATA_DIR
    prompt_clean = prompt.replace(" ", "_")
    Files = [f'{prompt_clean}{i}.jpg' for i in range(1, 5)]
    for jpg_file in Files:
        image_path = os.path.join(folder_path, jpg_file)
        try:
            img = Image.open(image_path)
            print(f'Opening image: {image_path}')
            img.show()
            await asyncio.sleep(1)
        except IOError:
            print(f'Unable to open image: {image_path}')
```



```
async def generate_images(prompt: str):
    """Generate 4 images using FLUX.1 [schnell]"""
    if client is None:
        print("[ImageGeneration] Error: HuggingFace client not initialized")
        return
    images = []
    for _ in range(4):
        try:
            image = await asyncio.wait_for(
                asyncio.to_thread(
                    client.text_to_image,
                    f'{prompt}, quality:4k, sharpness:maximum, Ultra High details, high resolution,
seed = {randint(0, 1000000)}',
                    model="black-forest-labs/FLUX.1-schnell",
                ),
                timeout=60 # Increased timeout for image generation
            )
            images.append(image)
        except asyncio.TimeoutError:
            print(f'[ImageGeneration] Timeout generating image')
        except Exception as e:
            print(f'[ImageGeneration] Error generating image: {e}')
    prompt_clean = prompt.replace(' ', '_')
    for i, img in enumerate(images):
        save_path = os.path.join(DATA_DIR, f'{prompt_clean}{i+1}.jpg')
        img.save(save_path)
        print(f'[ImageGeneration] Saved: {save_path}')
    async def GenerateImages(prompt: str):
        """Main function to generate and display images"""
        print(f'[ImageGeneration] Generating images for: {prompt}')
```

```
await generate_images(prompt)
await open_images(prompt)
```

5.5.4 Language Translation

File: [Translator.py]

```
import requests
from gtts import gTTS
import os
from playsound import playsound
import time
from speech import speak
def translategl(query):
    """Translate query using Google Translate API via requests library"""
    speak("SURE SIR")
    query = query.strip()
    print(f"[TRANSLATOR] Query: '{query}'")
    # Language mapping
    languages = {
        'af': 'Afrikaans', 'ar': 'Arabic', 'bg': 'Bulgarian', 'bn': 'Bengali',
        'ca': 'Catalan', 'cs': 'Czech', 'cy': 'Welsh', 'da': 'Danish',
        'de': 'German', 'el': 'Greek', 'en': 'English', 'es': 'Spanish',
        'et': 'Estonian', 'fa': 'Persian', 'fi': 'Finnish', 'fr': 'French',
        'gu': 'Gujarati', 'he': 'Hebrew', 'hi': 'Hindi', 'hr': 'Croatian',
        'hu': 'Hungarian', 'id': 'Indonesian', 'it': 'Italian', 'ja': 'Japanese',
        'kn': 'Kannada', 'ko': 'Korean', 'lt': 'Lithuanian', 'lv': 'Latvian',
        'mk': 'Macedonian', 'ml': 'Malayalam', 'mr': 'Marathi', 'ne': 'Nepali',
        'nl': 'Dutch', 'pa': 'Punjabi', 'pl': 'Polish', 'pt': 'Portuguese',
        'ro': 'Romanian', 'ru': 'Russian', 'sk': 'Slovak', 'sl': 'Slovenian',
        'so': 'Somali', 'sq': 'Albanian', 'sv': 'Swedish', 'ta': 'Tamil',
        'te': 'Telugu', 'th': 'Thai', 'tl': 'Tagalog', 'tr': 'Turkish',
```

```

        'uk': 'Ukrainian', 'ur': 'Urdu', 'vi': 'Vietnamese', 'zh-CN': 'Chinese (Simplified)',
        'zh-TW': 'Chinese (Traditional)'
    }

    name_to_code = {v.lower(): k for k, v in languages.items()}
    print("\n[TRANSLATOR] Available languages:")
    for code, name in sorted(languages.items()):
        print(f" {code}: {name}")
        speak("Choose the language in which you want to translate")
    user_input = input("To_Lang :- ").strip().lower()
    lang_code = name_to_code.get(user_input, user_input)
    print(f"[TRANSLATOR] Target language: {lang_code}")
    if lang_code not in languages:
        speak("Language not supported")
    return

```

5.5.5 Safety & Security

File: [Safety.py]

Safety.py - Content validation

import re

BLOCKED = {

"porn", "nude", "bomb", "hack", "drugs", "steal",

"kill myself", "suicide", "hate speech"

}

EDUCATIONAL = ["what is", "explain", "tell me", "learn", "how does"]

def check_safety(query):

"""Validate query for safety"""

q_lower = query.lower()

for word in BLOCKED:

if word in q_lower:

Allow if educational'

```
    if any(p in q_lower for p in EDUCATIONAL):
        return True, "Educational"
    return False, f"Blocked: {word}"
# Check for injection patterns
if re.search(r'[$()|s;]', query):
    return False, "Suspicious"
return True, "Safe"
```

5.5.6 Face Authentication

File: [FaceAuth.py]

FaceAuth.py - Biometric auth with liveness

```
import cv2, face_recognition, dlib, numpy as np
```

```
from scipy.spatial import distance as dist
```

```
def eye_aspect_ratio(eye):
```

```
    """Calculate EAR for blink detection"""
```

```
    A = dist.euclidean(eye[1], eye[5])
```

```
    B = dist.euclidean(eye[2], eye[4])
```

```
    C = dist.euclidean(eye[0], eye[3])
```

```
    return (A + B) / (2.0 * C)
```

```
def detect_blink(cam, detector, predictor, timeout=5):
```

```
    """Detect blink for liveness"""
```

```
    L_EYE = list(range(36, 42))
```

```
    R_EYE = list(range(42, 48))
```

```
    start = cv2.getTickCount() / cv2.getTickFrequency()
```

```
    consecutive_close = 0
```

```
    while (cv2.getTickCount() / cv2.getTickFrequency() - start) < timeout:
```

```
        ret, frame = cam.read()
```

```
        if not ret: break
```

```
            gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
```

```
            faces = detector(gray, 0)
```

```
        if len(faces) == 0: continue
        shape = predictor(gray, faces[0])
        pts = np.array([[p.x, p.y] for p in shape.parts()])
        l_ear = eye_aspect_ratio(pts[L_EYE])
        r_ear = eye_aspect_ratio(pts[R_EYE])
        avg = (l_ear + r_ear) / 2.0
        if avg < 0.25:
            consecutive_close += 1
        elif consecutive_close >= 2:
            return True # Blink detected
        return False
def authenticate():
    """Face auth with liveness check"""
    import pickle
    with open("Data/user_face.pkl", 'rb') as f:
        known = pickle.load(f)
        cam = cv2.VideoCapture(0)
        detector = dlib.get_frontal_face_detector()
        predictor = dlib.shape_predictor("Models/shape_predictor_68_face_landmarks.dat")
        if not detect_blink(cam, detector, predictor):
            print("Liveness check failed")
            return False
        ret, frame = cam.read()
        locs = face_recognition.face_locations(frame)
        if len(locs) == 0:
            return False
        live = face_recognition.face_encodings(frame, locs)[0]
        match = face_recognition.compare_faces([known], live, tolerance=0.6)
        cam.release()
    return match[0]
```

5.6 Screenshots


Command mode:



```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
[DEBUG] Welcome message completed
[AUTH] User authorized.
[STARTUP] Waiting for models to complete loading...
[STARTUP] Models ready!
[STARTUP] Starting automation scheduler...
[AUTOMATION] Scheduler started
WELCOME SIR! PLZ SPEAK [WAKE UP] TO LOAD ME UP
Listening (en-in)...
[DEBUG] Awake=False, Query='none'
Listening (en-in)...
```

Fig 1.8 (Command mode output)

Face authentication result:




```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
[DEBUG] Moving to face matching phase...
[SPEAK DEBUG] Attempting to speak: Liveness confirmed. Verifying identity.
[SPEAK DEBUG] Speech completed successfully

[Phase 2] Face Matching
[Success] Face authenticated!
[DEBUG] Preparing welcome message...
[DEBUG] Speaking: Verification successful. Welcome back sir. Good Afternoon, sir. Please tell me how can I help you?
[SPEAK DEBUG] Attempting to speak: Verification successful. Welcome back sir. Good Afternoon, sir. Please tell me how can I help you?
```

Fig 1.9 (Face Auth Output)

Wake word (Hotword) detection:



```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
WELCOME SIR! PLZ SPEAK [WAKE UP] TO LOAD ME UP
Listening (en-in)...
[Noise] Applied noise cancellation
Understanding...
You Said: wake up
[DEBUG] Awake=False, Query='wake up'
DEBUG: Matched words: ['wake', 'up']
DEBUG: Prob: wake_up -> 0.9989446997642517
[DEBUG] process_command_execution called with: 'wake up'
Outer Loop Prediction: wake_up (0.9989447)
Predicted Entity: wake_up (0.9989447)
```

Fig 2.0 (Hotword Detection Output)

Web Interface:




Fig 2.1 (Web Interface Output)

Image generation output:





Fig 2.2 (Image Generation Output)

Opening simultaneously YouTube and Notepad:






Fig 2.3 (Compound Command Execution)


Opening common desktop apps:




WhatsApp



Notepad



Chrome



Calculator

Fig 2.4 (Opening Desktop Apps Output)

Text responses through web interface:




Fig 2.5 (Text Response Output)

Playing music through YouTube:




Fig 2.6 (Playing music from YT Output)

Getting information from google:

```
Listening (en-in)...
[Noise] Applied noise cancellation
Understanding...
You Said: get information about Narendra Modi from Google
[DEBUG] Awake=False, Query='get information about narendra modi from google'
DEBUG: Matched words: ['get', 'information', 'about', 'from', 'google']
DEBUG: Prob: google_search -> 1.0
Listening (en-in)...
```




Fig 2.7 (Google Search Output – Information Retrieval)

Language translation:

```
[SPEAK DEBUG] Speech completed successfully
To_Lang :- hi
[TRANSLATOR] Target language: hi
[TRANSLATOR] Translating 'i want to eat pizza' to hi...
[TRANSLATOR] ✓ Translation: 'मैं पज़्ज़ा खाना चाहता हूँ'
[SPEAK DEBUG] Attempting to speak: मैं पज़्ज़ा खाना चाहता हूँ
[SPEAK DEBUG] Speech completed successfully
[TRANSLATOR] Playing audio (lang=hi)...
[TRANSLATOR] ✓ Audio played
```




Fig 2.8 (Language Translation Output)

Checking internet speeds:




Fig 2.9 (Getting Wifi Speeds Output)

CHAPTER – 6

TESTING

Testing is an essential phase in the development of the Jarvis AI system. It ensures that all modules of the assistant function correctly and interact properly with each other. Since the system integrates multiple technologies such as speech recognition, Natural Language Processing (NLP), machine learning models, automation modules, and conversational AI, thorough testing is required to verify system accuracy, performance, and reliability.

The Jarvis AI system was tested using different voice commands, automation scenarios, and conversational inputs to ensure that each component performs as expected. Both module-level testing and system-level testing were carried out during development.

The testing process focused on validating:

- speech recognition accuracy,
- intent classification correctness,
- automation task execution,
- conversational AI responses,
- translation functionality,
- and overall system performance.

6.1 Testing Strategy

The testing strategy followed a modular testing approach, where each module was tested individually before integrating them into the complete system.

The following testing methods were used:

- Unit Testing
- Integration Testing
- System Testing
- Functional Testing
- User Testing

6.1.1 Unit Testing

Each module of the Jarvis AI system was tested independently to ensure correct functionality.

Examples:

- Speech recognition module tested with different voice inputs.
- Intent classification tested with multiple command patterns.
- Automation module tested by executing system commands.
- Translation module tested with different languages.
- Conversational module tested with open-ended queries.

Unit testing ensured that each component performed correctly before integration.

6.1.2 Integration Testing

After unit testing, modules were combined and tested together to verify communication between components.

Integration testing included:

- Speech recognition → NLP processing
- NLP processing → Intent classification
- Intent classification → Automation module
- Conversational fallback → LLM module

This ensured smooth data flow between modules.

6.1.3 System Testing

System testing evaluated the Jarvis AI assistant as a complete system.

The assistant was tested for:

- voice command execution,
- automation tasks,
- translation features,
- conversational interaction,
- workflow execution.

System testing confirmed that the assistant operates correctly in real-time conditions.

6.1.4 Functional Testing

Functional testing verified whether the assistant performs the intended tasks.

Key features tested include:

- Opening applications using voice commands
- Searching information online
- Scheduling tasks
- Generating images
- Performing translation
- Responding to conversational queries

All functional requirements were validated successfully.

6.1.5 User Testing

User testing was performed by interacting with the assistant using real voice commands. This helped evaluate usability, response accuracy, and interaction quality.

The assistant responded correctly to most commands and demonstrated stable performance.

6.2 Test Cases

The following table shows sample test cases used to validate system functionality.

Test Case ID	Module	Input Command	Expected Output	Result
TC01	Hotword Detection	Hey Jarvis/ Wakeup	Assistant activates	Pass
TC02	Speech Recognition	What time is it?	Assistant speaks current time	Pass
TC03	Intent Classification	Search Google for AI	Browser opens with search results	Pass
TC04	Automation	Open Chrome	Chrome application opens	Pass
TC05	Automation	Close Notepad	Notepad closes	Pass
TC06	NLP Processing	How are you	Assistant responds conversationally	Pass
TC07	Translation	Translate hello to Hindi	Hindi translation spoken	Pass
TC08	Sentiment Analysis	I am very happy	Positive sentiment detected	Pass
TC09	Named Entity Recognition	Send message to Rahul	Name detected correctly	Pass
TC10	Scheduling	Set alarm for 7 AM	Alarm scheduled	Pass
TC11	Workmode	Run morning workflow	Workflow executed	Pass
TC12	Speech Mode	Switch to offline mode	Vosk STT activated	Pass
TC13	Automation	Open Chrome	Chrome application opens	Pass

TC14	Conversational AI	Tell me a joke	LLM response generated	Pass
TC15	Speech Mode	Switch to online mode	Google STT activated	Pass
TC16	Information Retrieval	Search Wikipedia for Python	Wikipedia page opens	Pass
TC17	Image Generation	Generate image of sunset	Image generated	Pass
TC18	Compound Command	Open YouTube and Notepad	Both apps open	Pass
TC19	System Control	Volume Up	System volume increases	Pass

Table 1.2 – Test Cases

6.3 Test Results

The testing results indicate that the Jarvis AI assistant performs reliably across different modules. Speech recognition worked effectively in both offline and online modes. The intent classification model correctly identified most user commands, and automation tasks were executed successfully.

The conversational AI module generated meaningful responses for open-ended queries, and the translation module performed accurately across multiple languages.

Minor variations in speech recognition accuracy were observed in noisy environments, but overall system performance remained stable.

6.4 Validation

Validation ensures that the developed system meets project objectives and user requirements.

The Jarvis AI assistant was validated based on:

- accuracy of intent classification,
- correctness of automation execution,
- quality of speech recognition,
- conversational response quality,
- system responsiveness.

The system successfully demonstrated:

- voice-based desktop automation,
- NLP-based command understanding,
- machine learning–based intent classification,
- conversational AI interaction,
- translation and image generation features.

Thus, the Jarvis AI system satisfies the functional requirements defined during system analysis.

CHAPTER – 7

RESULTS AND DISCUSSIONS

This chapter presents the results obtained from the implementation and testing of the Jarvis AI — Desktop Based Voice Assistant. The performance of the assistant is evaluated based on speech recognition accuracy, intent classification correctness, automation execution, conversational response quality, and overall system usability.

The Jarvis AI system successfully integrates speech recognition, NLP processing, machine learning models, automation modules, and conversational AI into a single intelligent assistant capable of performing desktop-level tasks using voice commands.

7.1 Output

The system was executed and tested using multiple voice commands covering automation, conversation, translation, scheduling, and image generation features.

The assistant demonstrated the following capabilities:

- Accurate speech-to-text conversion in both offline and online modes
- Correct intent classification for most commands
- Successful execution of automation tasks
- Intelligent responses from the conversational AI module
- Translation between multiple languages
- Image generation using generative AI models
- Workflow execution and scheduling

Screenshots of these outputs are included in the Implementation chapter.

7.2 Performance Analysis

The performance of the Jarvis AI assistant was evaluated based on system responsiveness, recognition accuracy, and execution reliability.

7.2.1 Speech Recognition Performance

The offline speech recognition module using Vosk provided fast response times and reliable performance without requiring internet connectivity. The online speech recognition module using Google Speech API provided higher transcription accuracy for complex commands.

The hybrid approach ensured both speed and accuracy.

7.2.2 Intent Classification Performance

The neural network-based intent classification model successfully identified user commands with high accuracy. The Bag-of-Words representation combined with lemmatization allowed the system to handle variations in user input.

The transformer-based classifier further improved language understanding for complex queries.

7.2.3 Automation Performance

The automation module executed system commands reliably, including:

- opening applications,
- performing searches,
- running workflows,
- scheduling tasks.

Execution time for automation tasks was minimal, and system operations were performed correctly.

7.2.4 Conversational AI Performance

The conversational AI module successfully handled queries that did not match predefined intents. The LLM integration generated meaningful responses and maintained conversational flow.

This significantly improved user interaction compared to rule-based assistants.

7.2.5 Translation and Image Generation Performance

The translation module produced accurate multilingual output using neural machine translation. The image generation module successfully generated images from text prompts using Stable Diffusion models.

Both features demonstrated the integration of advanced AI technologies into the assistant system.

7.3 Discussion

The results demonstrate that the Jarvis AI system successfully achieves its objective of providing a **voice-controlled desktop assistant capable of intelligent interaction and automation**.

The modular architecture allowed seamless integration of:

- speech recognition,
 - NLP processing,
 - machine learning models,
 - automation modules,
 - conversational AI,
 - translation,
 - generative AI features.
-
- ✓ The system performed reliably in both offline and online environments. Voice interaction reduced manual effort and improved user productivity.
 - ✓ Although minor variations in speech recognition accuracy were observed in noisy environments, overall system performance remained stable and responsive.
 - ✓ The integration of multiple AI technologies in a single assistant system demonstrates the practical application of machine learning and NLP techniques in desktop automation.

CHAPTER – 8

CONCLUSION AND FUTURE SCOPE

8.1 Conclusion

The Jarvis AI — Desktop Based Voice Assistant project demonstrates the successful integration of Artificial Intelligence, Natural Language Processing (NLP), Machine Learning (ML), speech recognition, and automation technologies into a single intelligent assistant system.

The developed assistant is capable of understanding voice commands, identifying user intent, executing desktop-level tasks, and generating conversational responses. The system supports both offline and online speech recognition, ensuring reliability and flexibility in different environments. The integration of machine learning–based intent classification, NLP processing, and automation modules enables the assistant to perform tasks efficiently and interact naturally with the user.

Additional features such as translation, workflow automation, conversational AI integration, image generation, and face tracking demonstrate the extensibility of the system. The modular architecture allows new functionalities to be added without affecting existing modules.

The project successfully meets its objectives by providing a voice-controlled desktop automation system that improves user productivity and interaction with the computer. It also demonstrates how AI technologies can be combined to create intelligent assistant applications.

Overall, the Jarvis AI system serves as a practical implementation of AI-based human–computer interaction using voice commands and automation.

8.2 Limitations

Although the Jarvis AI assistant performs effectively, some limitations exist.

Speech recognition accuracy may decrease in noisy environments or with unclear pronunciation. The assistant currently relies on predefined intent datasets, which may limit understanding of completely new commands. Some advanced features, such as LLM-based conversation and image generation, require internet connectivity.

The assistant is designed primarily for desktop environments, and performance may vary depending on system hardware and microphone quality. Additionally, the system does not yet support continuous learning from user behaviour.

8.3 Future Scope

The Jarvis AI assistant can be further improved by adding new features and capabilities.

Future enhancements may include:

- **Multilingual voice recognition support** for additional languages and accents.
- **Adaptive learning mechanisms** to personalize responses based on user behaviour.
- **Improved noise-robust speech recognition** for better performance in real-world environments.
- **Mobile and web deployment** of the assistant system.
- **Integration with IoT** for extended automation.
- **Enhanced security features** such as biometric authentication and encrypted data storage.
- **More advanced conversational AI models** for deeper context understanding.
- **Plugin-based architecture** for adding new assistant skills easily.

These improvements would make Jarvis AI more intelligent, scalable, and suitable for real-world applications.

CHAPTER - 9

REFERENCES

- [1] R. Gudla, P. M. Kumar, F. Shaik, and A. V. Pilla, “AI Voice Assistant using NLP and Python Libraries (M.A.R.F),” *International Journal of Innovative Science and Research Technology*, vol. 9, no. 2, 2024.
- [2] B. S. S. Shabana, S. D. Abarajitha, and M. K. Taufeeq, “An AI Powered Voice Assistant for Enhanced User Interaction,” in *Proc. IEEE International Conference on Communication and Signal Processing (ICCSP)*, 2025.
- [3] A. Gobinath, P. Prakash, C. Manjula Devi, et al., “Voice Assistant with AI Chat Integration using OpenAI,” in *Proc. International Conference on Intelligent Techniques in Control, Optimization and Signal Processing (INCOS)*, 2024.
- [4] L. Lazzaroni, F. Bellotti, and R. Berta, “An Embedded End-to-End Voice Assistant,” *Engineering Applications of Artificial Intelligence*, Elsevier, 2024.
- [5] V. Jain, Y. Patidar, U. Jain, V. Parwal, A. Kumawat, and S. Sureliya, “Automate Personal Voice Based Assistant Using Python,” in *Proc. Asian Conference on Innovation in Technology (ASIANCON)*, IEEE, 2024.
- [6] Python Software Foundation, “Python Documentation.” [Online]. Available: <https://docs.python.org>
- [7] TensorFlow, “TensorFlow Documentation.” [Online]. Available: <https://www.tensorflow.org>
- [8] Hugging Face, “Transformers Documentation.” [Online]. Available: <https://huggingface.co/docs>
- [9] Alpha Cephei, “Vosk Speech Recognition Toolkit.” [Online]. Available: <https://alphacephei.com/vosk>
- [10] Google Cloud, “Speech-to-Text API Documentation.” [Online]. Available: <https://cloud.google.com/speech-to-text>
- [11] “Groq API Documentation.” [Online]. Available: <https://console.groq.com/docs>
- [12] Stability AI, “Stable Diffusion Documentation.” [Online]. Available: <https://stability.ai>